

Data Exploration, Visualization, and Feature Engineering

Data Science Dojo

Agenda

- Why data exploration and visualization
- Exploration and visualization using R
 - Core R functionality – iris dataset
 - lattice package – mtcars dataset
 - ggplot2 package – diamonds and G20 datasets
- Story-telling with data
 - Titanic data set

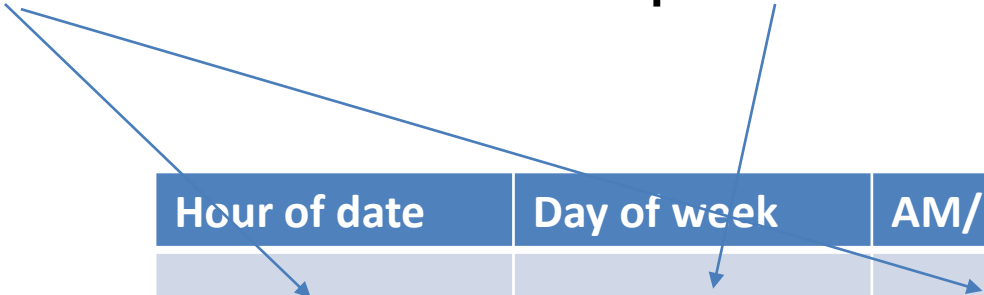
WHY DATA EXPLORATION AND VISUALIZATION

Data beats algorithm but...

- More data usually yields good generalization performance, even with a simple algorithm
- But there are caveats
 - Amount of data may have diminishing returns
 - Data quality and variety matters
 - A decent performing learning algorithm is still needed
 - Most importantly, extracting useful features out of data is important

Why feature engineering matters

- 23:05:33 –5 UTC, April 3, 2014



| Hour of date | Day of week | AM/PM |
|--------------|-------------|-------|
| | | |
| | | |

Dispelling common myths

- There is *NO* single ML algorithm that will take raw data and give you the best model



- You do *NOT* need to know a lot of machine learning algorithms to build robust predictive models

Janitorial work is important

- Not spending time on understanding your data is a source of many problems!
- Remember the 80/20 rule
 - 80% : Data cleaning, data exploration, feature engineering, pre-processing etc...
 - 20% : Model building

EXPLORATION AND VISUALIZATION USING R





Objectives

- Develop an understanding of the high-level thinking process of data exploration
- Make sense of data using visualization techniques
- Learn to perform feature engineering
- Become a good storyteller

Anscombe's quartet

| I | | II | | III | | IV | |
|------|-------|------|------|------|-------|------|-------|
| x | y | x | y | x | y | x | y |
| 10.0 | 8.04 | 10.0 | 9.14 | 10.0 | 7.46 | 8.0 | 6.58 |
| 8.0 | 6.95 | 8.0 | 8.14 | 8.0 | 6.77 | 8.0 | 5.76 |
| 13.0 | 7.58 | 13.0 | 8.74 | 13.0 | 12.74 | 8.0 | 7.71 |
| 9.0 | 8.81 | 9.0 | 8.77 | 9.0 | 7.11 | 8.0 | 8.84 |
| 11.0 | 8.33 | 11.0 | 9.26 | 11.0 | 7.81 | 8.0 | 8.47 |
| 14.0 | 9.96 | 14.0 | 8.10 | 14.0 | 8.84 | 8.0 | 7.04 |
| 6.0 | 7.24 | 6.0 | 6.13 | 6.0 | 6.08 | 8.0 | 5.25 |
| 4.0 | 4.26 | 4.0 | 3.10 | 4.0 | 5.39 | 19.0 | 12.50 |
| 12.0 | 10.84 | 12.0 | 9.13 | 12.0 | 8.15 | 8.0 | 5.56 |
| 7.0 | 4.82 | 7.0 | 7.26 | 7.0 | 6.42 | 8.0 | 7.91 |
| 5.0 | 5.68 | 5.0 | 4.74 | 5.0 | 5.73 | 8.0 | 6.89 |

Plot



Anscombe's quartet

| I | | II | | III | | IV | |
|------|-------|------|------|------|-------|------|-------|
| x | y | x | y | x | y | x | y |
| 10.0 | 8.04 | 10.0 | 9.14 | 10.0 | 7.46 | 8.0 | 6.58 |
| 8.0 | 6.95 | 8.0 | 8.14 | 8.0 | 6.77 | 8.0 | 5.76 |
| 13.0 | 7.58 | 13.0 | 8.74 | 13.0 | 12.74 | 8.0 | 7.71 |
| 9.0 | 8.81 | 9.0 | 8.77 | 9.0 | 7.11 | 8.0 | 8.84 |
| 11.0 | 8.33 | 11.0 | 9.26 | 11.0 | 7.81 | 8.0 | 8.47 |
| 14.0 | 9.96 | 14.0 | 8.10 | 14.0 | 8.84 | 8.0 | 7.04 |
| 6.0 | 7.24 | 6.0 | 6.13 | 6.0 | 6.08 | 8.0 | 5.25 |
| 4.0 | 4.26 | 4.0 | 3.10 | 4.0 | 5.39 | 19.0 | 12.50 |
| 12.0 | 10.84 | 12.0 | 9.13 | 12.0 | 8.15 | 8.0 | 5.56 |
| 7.0 | 4.82 | 7.0 | 7.26 | 7.0 | 6.42 | 8.0 | 7.91 |
| 5.0 | 5.68 | 5.0 | 4.74 | 5.0 | 5.73 | 8.0 | 6.89 |

Consider the 4 following different datasets

| | |
|---------------------------|-------|
| Mean of X | 9 |
| Variance of X | 11 |
| Mean of Y | 7.5 |
| Variance of Y | 4.125 |
| Correlation between X & Y | 0.816 |

Awareness



New to R?

- Focus on ideas/concepts rather than exact syntax. R help is your friend. 😊
 - ?mean, ?sd
 - ??melt (*use two question marks for packages not loaded*)
 - help()
 - example()
- All slides have code samples
- Sample code + slides: 'Data Exploration and Visualization' folder

Common graphical parameters

- Title of graph using the **main** function, `main = "title"`
- Label x- axis by using the **xlab** function, `xlab = "label x axis"`
- Label y- axis by using the **ylab** function, `ylab = "label y axis"`
- Colors controlled by **col**
- Get legends of layered plots with **auto.key=TRUE**

Exploring data commands

| Commands | Description |
|----------------------------------|---|
| read.csv() , read.table() | Load data/file into a dataframe |
| data() | Loads or resets a dataset |
| names() | List names of variables in a dataframe |
| head() | First 6 rows of data |
| tail() | Last 6 rows of data |
| str() | Display internal structure if R object |
| View() | View dataset in spreadsheet format in RStudio |
| dim() | Dimensions(rows and columns) of dataframe |
| summary() | Display 5-number summary and mean |
| colnames() | Provide column names |

CORE R GRAPHICS

The iris dataset


```
data(iris)
head(iris)
```

```
> head(iris)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4          0.2  setosa
2          4.9         3.0          1.4          0.2  setosa
3          4.7         3.2          1.3          0.2  setosa
4          4.6         3.1          1.5          0.2  setosa
5          5.0         3.6          1.4          0.2  setosa
6          5.4         3.9          1.7          0.4  setosa
```

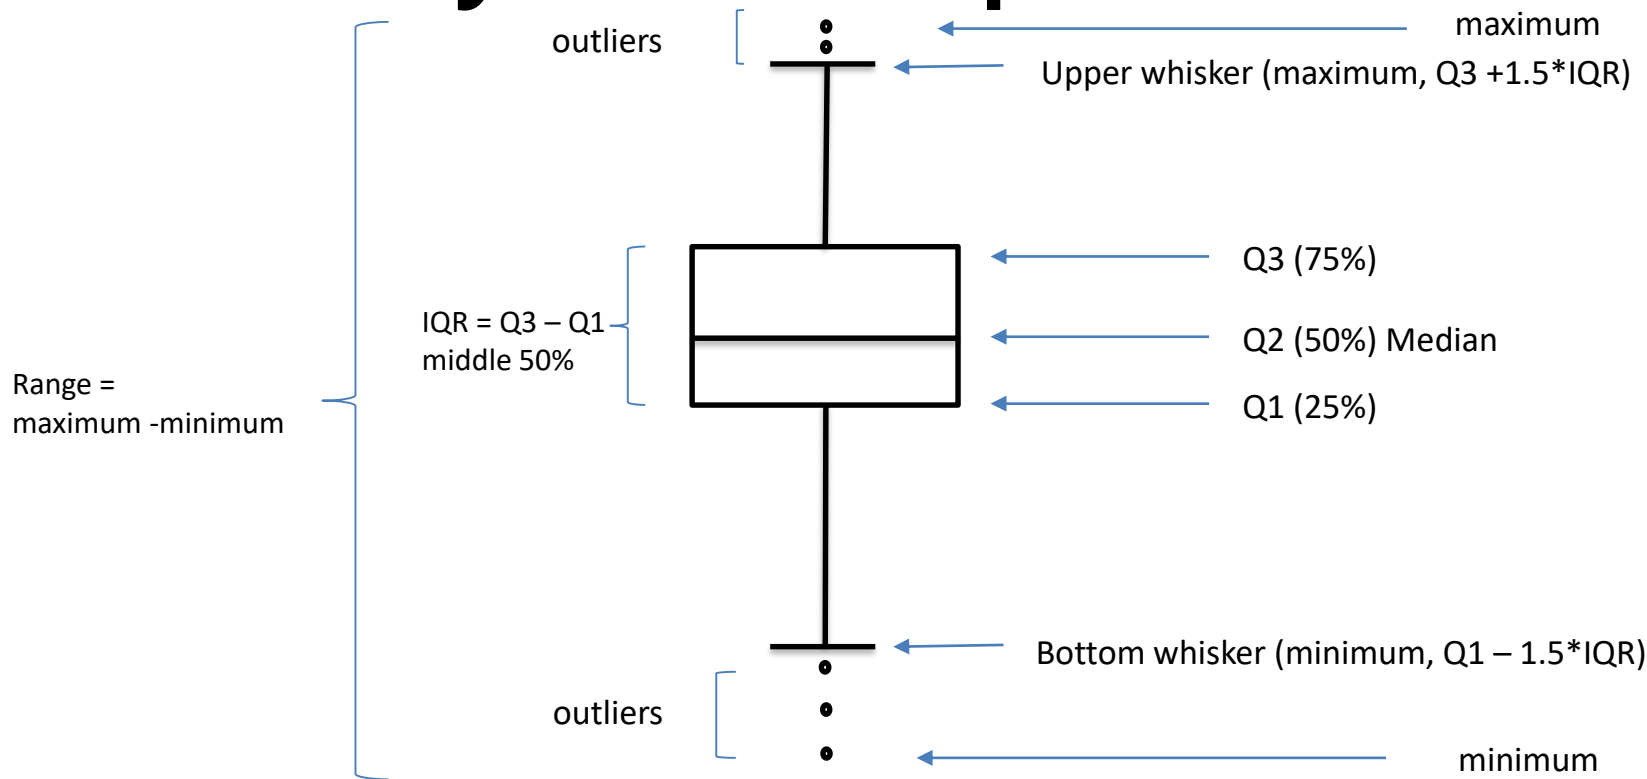
Boxplots

- Summarizes *quantitative/numeric* data

```
# Core Graphics
boxplot(
  Sepal.Length ~ Species,
  data=iris,
  main="Sepal Length for
  Various Species",
  xlab="Species",
  ylab="Sepal Length"
)
```




Anatomy of a boxplot



Plot

- Visual depiction of correlation between numeric variables

```
# Core Graphics
plot(Sepal.Length ~ Sepal.Width,
data=iris, xlab= "Sepal Length",
ylab= "Sepal Width")
```




LATTICE GRAPHICS


xyplot

- Plot counterpart in lattice package.
- Similar output as core graphics, but easier to color and segment points

```
# Lattice Graphics  
library(lattice)  
xyplot(Sepal.Width ~  
Sepal.Length, data=iris,  
groups=Species,  
auto.key=TRUE  
)
```



Density plots



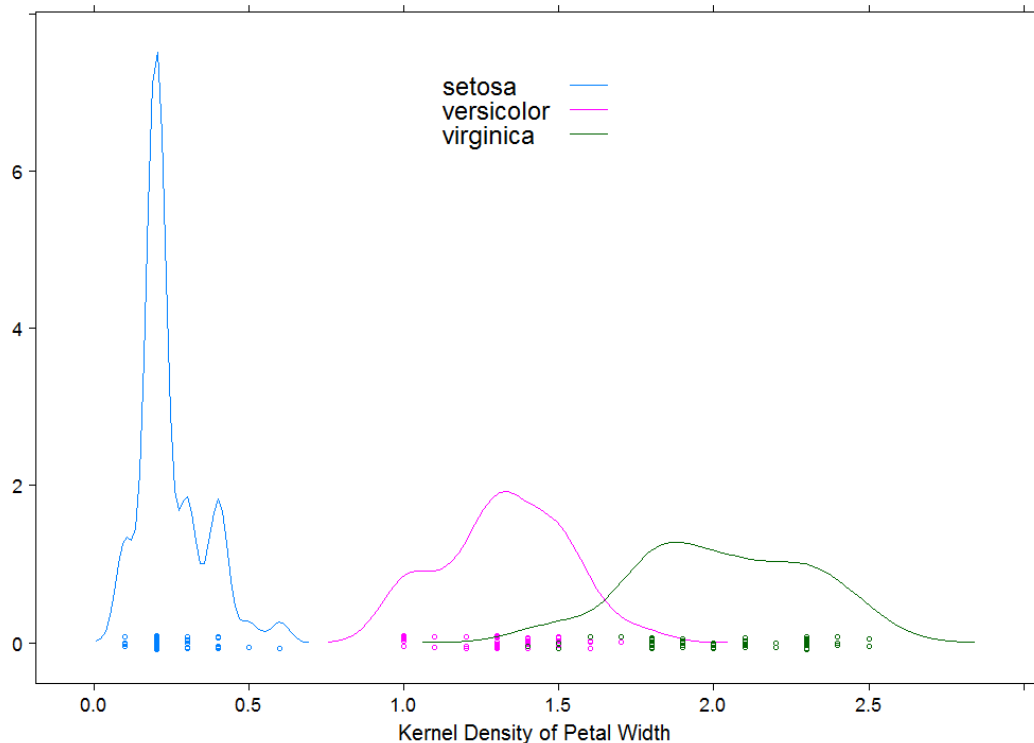
- Estimates density function from counts
- Area under the curve is always one
- Does not work with missing values

```
densityplot(iris$Petal.Length,  
            main="Kernel Density of  
            Petal Length", xlab="Petal  
            Length")
```

Try adding `plot.points=F`

Multiple density plots

Density of Petal Width by Species




```
densityplot(~Petal.Width,  
data=iris,  
groups=Species,  
auto.key=TRUE,  
xlab="Kernel Density of  
Petal Width",  
ylab="Frequency",  
main=list(label="Density  
of Petal Width by  
Species"))
```


Scatterplot matrix


- Multiple relationships in one graph
- Good for initial explorations

```
# Core Graphics
```

```
pairs(  
iris[,1:4],  
main="Scatterplot Matrix"  
)
```



Scatterplot matrix



```
# Lattice Graphics
```

```
splom(iris[1:4],  
groups=iris$Species)
```

In-class Exercise

- Using the "mtcars" dataset, predict mpg based on other columns.
- Create at least 2 different plots illustrating useful relationships in data and summarize your findings.

The “mtcars” dataset

```
data(mtcars)
head(mtcars)
```

```
> head(mtcars)
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|-------------------|------|-----|------|-----|------|-------|-------|----|----|------|------|
| Mazda RX4 | 21.0 | 6 | 160 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 |
| Mazda RX4 Wag | 21.0 | 6 | 160 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 |
| Datsun 710 | 22.8 | 4 | 108 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 |
| Hornet 4 Drive | 21.4 | 6 | 258 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 |
| Hornet Sportabout | 18.7 | 8 | 360 | 175 | 3.15 | 3.440 | 17.02 | 0 | 0 | 3 | 2 |
| Valiant | 18.1 | 6 | 225 | 105 | 2.76 | 3.460 | 20.22 | 1 | 0 | 3 | 1 |

GGPLOT2 GRAPHICS

The “diamonds” dataset

```
library(ggplot2)
data(diamonds)
head(diamonds)
```

```
> head(diamonds)
```

| | carat | cut | color | clarity | depth | table | price | x | y | z |
|---|-------|-----------|-------|---------|-------|-------|-------|------|------|------|
| 1 | 0.23 | Ideal | E | SI2 | 61.5 | 55 | 326 | 3.95 | 3.98 | 2.43 |
| 2 | 0.21 | Premium | E | SI1 | 59.8 | 61 | 326 | 3.89 | 3.84 | 2.31 |
| 3 | 0.23 | Good | E | VS1 | 56.9 | 65 | 327 | 4.05 | 4.07 | 2.31 |
| 4 | 0.29 | Premium | I | VS2 | 62.4 | 58 | 334 | 4.20 | 4.23 | 2.63 |
| 5 | 0.31 | Good | J | SI2 | 63.3 | 58 | 335 | 4.34 | 4.35 | 2.75 |
| 6 | 0.24 | Very Good | J | VVS2 | 62.8 | 57 | 336 | 3.94 | 3.96 | 2.48 |

ggplot fundamentals

- `ggplot()` provides a blank canvas for plotting
- `geom_*()` creates actual graphical layers
 - `geom_point()`
 - `geom_boxplot()`
- `aes()` defines an "aesthetic" either globally or by layer

Layering

```
ggplot(diamonds, aes()) + geom_point()
```




Layer 1



Layer 2


Density plot



```
ggplot(diamonds) +  
  geom_density(aes(x=carat),  
    fill="gray50")
```


- Note the location of aes()

Scatterplot



```
ggplot(diamonds,  
  aes(x=carat,y=price)) +  
  geom_point()
```

ggplot object



color

D
E
F
G
H
I
J


```
# Store the plot for future  
modification
```

```
g <- ggplot(diamonds,  
aes(x=carat, y=price))
```

```
# add settings specific to  
geom_point layer
```

```
g + geom_point(aes(color=color))
```

ggplot object



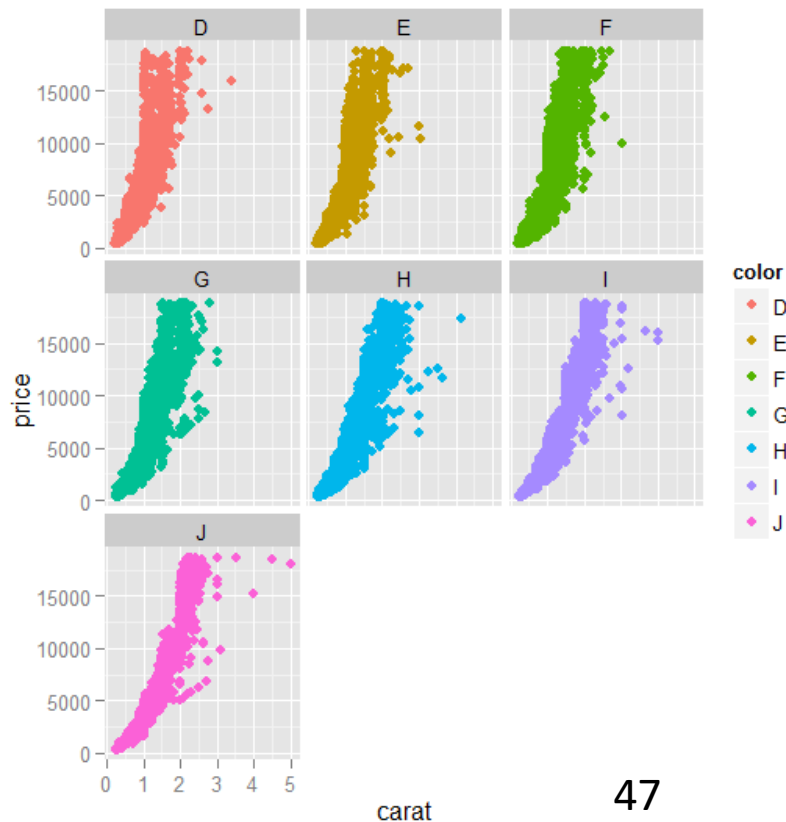
```
# Store the plot for future  
modification
```

```
g <- ggplot(diamonds,  
aes(x=carat, y=price))
```

```
# add settings specific to  
geom_point layer
```

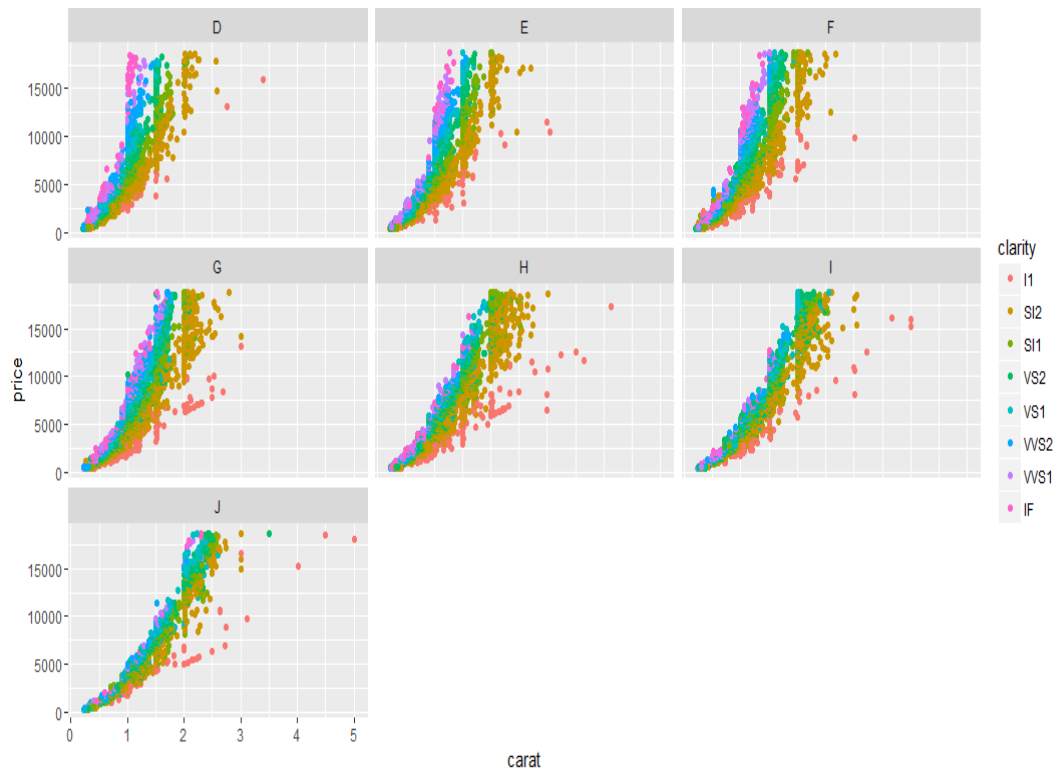
```
g +  
geom_point(aes(color=clarity))
```

Separating segments




```
g +  
  geom_point(aes(color=color)) +  
  facet_wrap(~ color)
```

Separating segments



```
g +  
  geom_point(aes(color  
=clarity)) +  
  facet_wrap(~ clarity)
```

More segments!

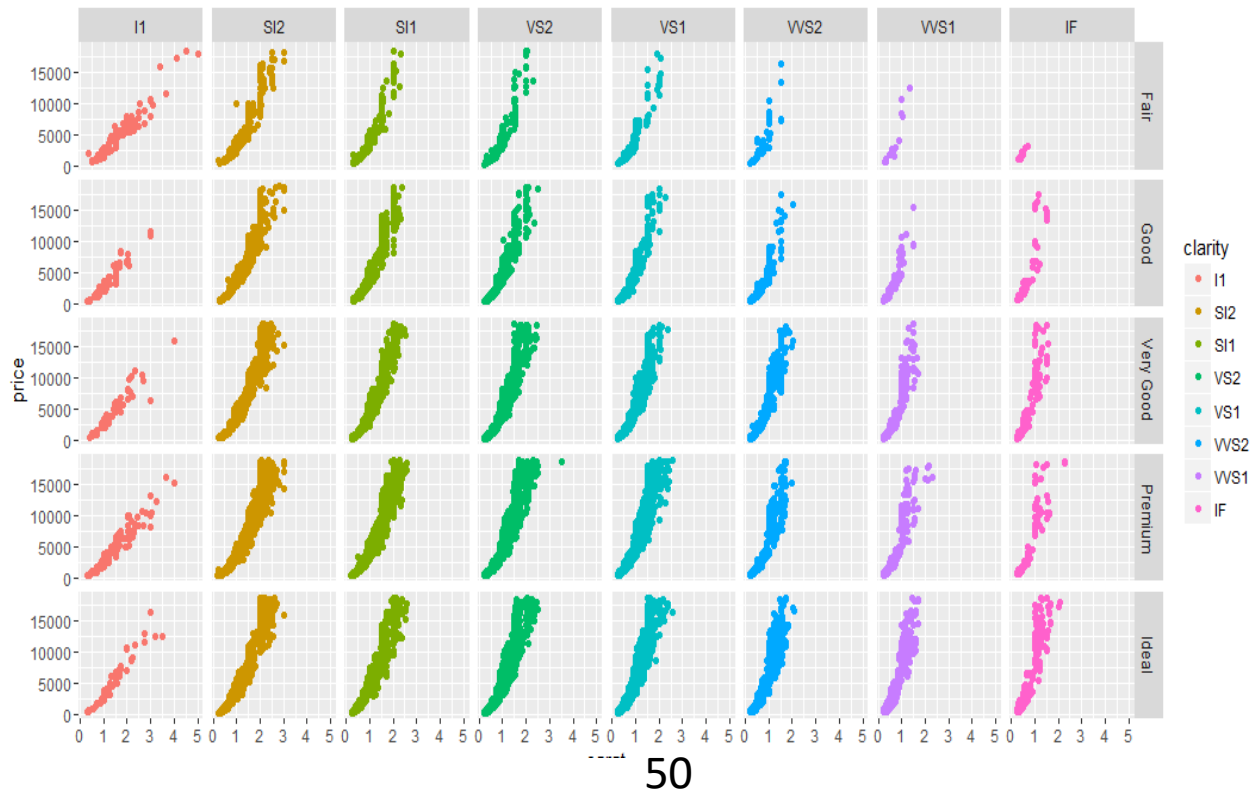


```
g +  
  geom_point(aes(color=color)) +  
  facet_grid(cut ~ clarity)
```

color

D
E
F
G
H
I
J

More segments!




```
g +  
  geom_point(aes(  
    color=clarity  
  )) +  
  facet_grid(cut  
    ~ clarity)
```


Heatmaps

- Helpful in understanding strength of relationship in a correlation or a distance matrix.
- Visualizes numbers by changing into colored cells.

Heatmaps




```
nba <-  
read.csv("http://datasets.flowingdata.com/  
ppg2008.csv", sep=";")
```

```
row.names(nba) <- nba$Name
```

```
nba_matrix <- data.matrix(nba)  
nba_heatmap <- heatmap(nba_matrix,  
Rowv=NA, Colv=NA, col = cm.colors(256),  
scale="column", margins=c(5,10))
```

Source: <http://flowingdata.com>

Heatmaps



```
nba <-  
read.csv("http://datasets.flowingdata.com/  
ppg2008.csv", sep=";")
```

```
row.names(nba) <- nba$Name
```

```
nba_heatmap <- heatmap(nba_matrix,  
Rowv=NA, Colv=NA, col =  
heat.colors(256), scale="column",  
margins=c(5,10))
```

Source: <http://flowingdata.com>

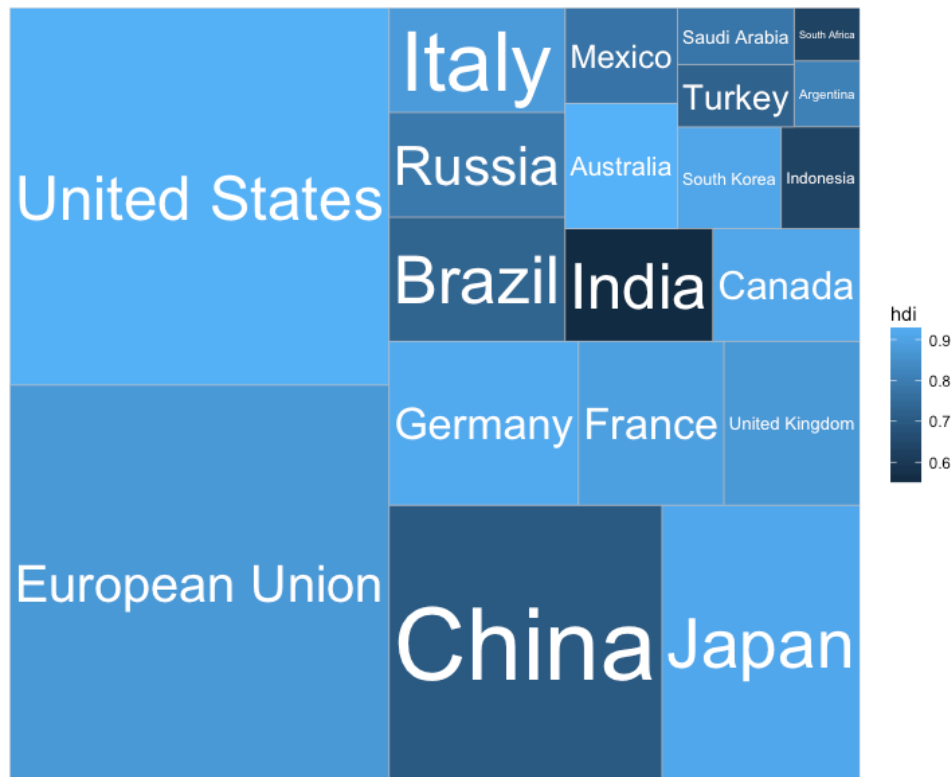
Treemap

- Powerful visualization of multi-dimensional values grouped by nested segments
- Provides an easy at-first-glance overview of a dataset's hierarchical structure
- Must specify variables for: area, fill, and label

Treemap

```
library(treemapify)
```

```
head(G20)
ggplot(G20, aes(area =
  gdp_mil_usd, fill = hdi,
  label = country)) +
  geom_treemap() +
  geom_treemap_text(color='white',
    place='center',grow=T)
```



Summary

- ✓ Basics of R
- ✓ Graphing in R – core, lattice, ggplot2, and treemapify
- ✓ Look at multiple types of graphs
- ✓ Visualize and segment data to gain more insights
- ✓ Identify key features
- ✓ Summarize findings

STORYTELLING WITH TITANIC

Finding the data set

- Set your working directory to the bootcamp root
- Load data in from "Datasets/titanic.csv"

Looking at the first few rows

```
titanic <- read.csv("titanic.csv") `
head(titanic)
```

```
> head(titanic)
  PassengerId Survived Pclass                    Name Sex Age SibSp Parch    Ticket   Fare Cabin Embarked
1          1         0       3        Braund, Mr. Owen Harris   male  22     1     0      A/5 21171   7.2500           S
2          2         1       1 Cumings, Mrs. John Bradley (Florence Briggs Thayer) female  38     1     0      PC 17599  71.2833   C85       C
3          3         1       3        Heikkinen, Miss. Laina female  26     0     0 STON/O2. 3101282   7.9250           S
4          4         1       1 Futrelle, Mrs. Jacques Heath (Lily May Peel) female  35     1     0      113803  53.1000  C123       S
5          5         0       3        Allen, Mr. William Henry   male  35     0     0      373450   8.0500           S
6          6         0       3        Moran, Mr. James     male  NA     0     0      330877   8.4583           Q
> |
```

What features should we consider?

What is the data type of each column?

```
str(titanic)
```

```
'data.frame':      891 obs. of  12 variables:
```

```
$ PassengerId: int  1 2 3 4 5 6 7 8 9 10 ...
```

```
$ Survived   : int  0 1 1 1 0 0 0 0 1 1 ...
```

```
$ Pclass     : int  3 1 3 1 3 3 1 3 3 2 ...
```

```
$ Name       : Factor w/ 891 levels "Abbing, Mr. Anthony",...: 109 191 358 277 16 559 520 629 417 581 ...
```

```
$ Sex        : Factor w/ 2 levels "female","male": 2 1 1 1 2 2 2 2 1 1 ...
```

```
$ Age        : num  22 38 26 35 35 NA 54 2 27 14 ...
```

```
$ SibSp      : int  1 1 0 1 0 0 0 3 0 1 ...
```

```
$ Parch      : int  0 0 0 0 0 0 0 1 2 0 ...
```

```
$ Ticket     : Factor w/ 681 levels "110152","110413",...: 524 597 670 50 473 276 86 396 345 133 ...
```

```
$ Fare       : num  7.25 71.28 7.92 53.1 8.05 ...
```

```
$ Cabin      : Factor w/ 148 levels "", "A10", "A14",...: 1 83 1 57 1 1 131 1 1 1 ...
```

```
$ Embarked   : Factor w/ 4 levels "", "C", "Q", "S": 4 2 4 4 4 3 4 4 4 2 ...
```

Casting

Set target column as a factor

```
titanic$Survived <- as.factor(titanic$Survived)
```


Rename factors and columns

```
levels(titanic$Survived) <- c("Dead", "Survived")  
levels(titanic$Embarked) <- c("Unknown", "Cherbourg",  
                             "Queenstown", "Southampton")  
str(titanic[,c("Embarked", "Survived")])
```

```
'data.frame':  891 obs. of  2 variables:  
 $ Embarked: Factor w/ 4 levels  
 "Unknown", "Cherbourg", ...: 4 2 4 4 4 3 4 ...  
 $ Survived: Factor w/ 2 levels "0","1": 1 2 2 2  
 1 1 1 1 2 2 ...
```

Class distribution: Pie Chart

```
survivedTable <- table(titanic$Survived)  
pie(survivedTable, labels=c("Died", "Survived"))
```



Is Sex a good predictor?

```
#Identify where sex = male for all columns
male <- titanic[titanic$Sex == "male",]


#Identify where sex = female for all columns
female <- titanic[titanic$Sex == "female",]

par(mfrow=c(1,2)) #two figures arranged in 1 row and 2
columns
pie(table(male$Survived), labels=c("Dead","Survived"))


pie(table(female$Survived), labels=c("Dead","Survived"))
```

Is **Sex** a good predictor?

Survival Proportion Among Men



Survival Proportion Among Women



Is Age a good predictor?

```
summary(titanic$Age)
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. | NA's |
|------|---------|--------|-------|---------|-------|------|
| 0.42 | 20.12 | 28.00 | 29.70 | 38.00 | 80.00 | 177 |

■ How about by Survival?


```
summary(titanic[titanic$Survived  
=="Dead",]$Age)
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. | NA's |
|------|---------|--------|-------|---------|-------|------|
| 1.00 | 21.00 | 28.00 | 30.63 | 39.00 | 74.00 | 125 |

```
summary(titanic[titanic$Survived  
=="Survived",]$Age)
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. | NA's |
|------|---------|--------|-------|---------|-------|------|
| 0.42 | 19.00 | 28.00 | 28.34 | 36.00 | 80.00 | 52 |


Sample solution



67

```
density(titanic$Age)
#NAs prevent this
> d <-
density(na.omit(titanic$Age))
> plot(d, main="Kernel
Density of Age")
>
polygon(d, border="green", l
wd=3)
```



Is Age a good predictor for Survival?



```
densityplot(~Age, data=titanic,  
groups=Survived, plot.points=F,  
lwd=3)
```

Note: won't work with missing values

Is Age a good predictor for Gender?



```
densityplot(~ Age,  
data=titanic, groups=Sex,  
plot.points=F, lwd=3)
```

QUESTIONS